

EE5907R Programming Assignment

20% of Final Grade

Project Deadline: 11.59pm, Monday, Feb 27, 2017

Submit the following electronic files in one zipped folder onto the “CA1 Submission” folder on IVLE. The zip filename **MUST** be "[name_on_matric_card]_[matric_number]_CA1.zip"

1. **The pdf file of a well-written, concise project report. The report should NOT be longer than 10 pages (font 12, single space, arial). The report filename MUST be "[name_on_matric_card]_[matric_number]_report.pdf".**
2. **Your source code folder.**
3. **Readme file containing instructions to run your code. This readme file MUST be inside your source code**
4. **Please don't include project data in your zip folder.**

Before you start, take note of the following:

1. You may discuss the assignment with your classmates, but must write the code completely on your own. Plagiarism will be severely punished.
2. The GA prefers you to use matlab for this assignment. If you are unable to use matlab, you can use octave or python.
3. The data (**spamData.mat**) can be downloaded from the IVLE workbin. The data is in matlab format, which can be quite easily read inside python with the right package. If you can't figure out the right package, you probably should be using matlab or octave :)
4. For all the questions, there are publicly available software libraries that implement some versions of these classifiers. However, implementing the algorithms yourself will help you understand the theory better. Therefore you are expected to implement the classifiers yourself.
5. The evaluation criteria include organization and clarity of the report, correctness of the implementation, and performance of the classifiers.
6. Please be considerate to your GA. Be as clear as possible in your submission, e.g., having a clear readme and provide helpful comments in your code. You are more likely to get a better grade if your GA can understand your code and report!

Acknowledgement: This assignment is a variation of a problem from Kevin Murphy's "Machine Learning: A Probabilistic Perspective" textbook.

Data Description

The data is an email spam dataset, consisting of 4601 email messages with 57 features. Feature descriptions are found here (<http://statweb.stanford.edu/~tibs/ElemStatLearn/datasets/spam.info.txt>). We have divided the data into a training set (3065 emails) and test set (1536 emails) with accompanying labels (1 = spam , 0 = not spam).

Data Processing

One can try different preprocessing of the features. Consider the following separately:

- (a) z-normalization: standardize each column so they have 0 mean and unit variance. Use empirical mean and empirical variance.
- (b) log-transform: transform each feature using $\log(x_{ij} + 0.1)$ (assume natural log)
- (c) binarization: binarize features: $\mathbb{I}(x_{ij} > 0)$, i.e., features that are equal to 0 remain as 0. Features that are greater than 0 are set to 1.

Q1. Beta-bernoulli Naive Bayes (24%)

Fit a Beta-bernoulli naive Bayes classifier on the binarized data from the Data Processing section. Since there are a lot of spam and non-spam emails, you do not need to assume any prior on the class label. In other words, the class label prior π_c can be estimated using ML and you can use π_c^{ML} as a plug-in estimator for testing.

On the other hand, you should assume a prior Beta(α, α) on the feature distribution (note that the two hyperparameters for the Beta prior are set to be equal). For each value of $\alpha = \{0, 0.5, 1, 1.5, 2, \dots, 100\}$, fit the classifier on the training data and compute its error rate (i.e., **percentage of emails classified wrongly**) on the test data.

Please use Bayesian (i.e., posterior predictive) training and testing (see **week 2** lecture notes on “Bayesian Training of Naive Bayes” and “Bayesian Testing of Naive Bayes”). Make sure you include at least the following in your report:

- Plots of training and test error rates versus α
- What do you observe about the training and test errors as α change?
- Training and testing error rates for $\alpha = 1, 10$ and 100.

Q2. Gaussian Naive Bayes (24%)

Fit a Gaussian naive Bayes classifier on the continuous data (both z-normalization and log-transform) from the Data Processing section. Since there are a lot of spam and non-spam emails, you do not need to assume any prior on the class label. In other words, the class label prior π_c can be estimated using ML and you can use π_c^{ML} as a plug-in estimator for testing.

For this exercise, just use maximum likelihood to estimate the class conditional mean and variance of each feature and use ML estimates as a plug-in estimator for testing (see **week 2** lecture notes on “Plug-in Testing of Naive Bayes” and **week 3** lecture notes on “ML estimation of μ, σ^2 ”). Make sure you include the following in your report:

- Training and testing error rates for both z-normalized and log-transformed data.

Q3. Logistic regression (24%)

For each version of data (z-normalization, log-transform and binarization) from the Data Processing section, fit a logistic regression model with l_2 regularization (see **week 4** lecture notes on “Newton’s Method for Logistic Regression” and “ l_2 regularization”). For each regularization parameter value $\lambda = \{1, 2, \dots, 9, 10, 15, 20, \dots, 95, 100\}$ (note the jump in interval from 10 to 15 and beyond), fit the logistic regression model on the training data and compute its error rate (i.e., **percentage of emails classified wrongly**) on the test data. Make sure you include at least the following in your report:

- Plots of training and test error rates versus λ
- What do you observe about the training and test errors as λ change?
- What do you observe about the error rates of the different preprocessing strategies?
- Training and testing error rates for $\lambda = 1, 10$ and 100 .

Don’t forget to include the bias term in the logistic regression and your l_2 regularization should not apply to the bias term.

Q4. K-Nearest Neighbors (24%)

For each version of the data (z-normalization, log-transform and binarization) from the Data Processing section, implement a KNN classifier (see **week 4** lecture notes on “Alternative Classification Approach”).

For each value of $K = \{1, 2, \dots, 9, 10, 15, 20, \dots, 95, 100\}$ (note the jump in interval from 10 to 15 and beyond), compute the training and test error rates (i.e., **percentage of emails classified wrongly**).

For the continuous data (z-normalization, log-transform) use the Euclidean distance to measure distance between neighbors. For the binarized data, use the Hamming distance (i.e., distance is equal to number of features that are different; https://en.wikipedia.org/wiki/Hamming_distance). Make sure you include at least the following in your report:

- Plots of training and test error rates versus K
- What do you observe about the training and test errors as K change? Why is training error not 0 when $K = 1$?
- What do you observe about the error rates of the different preprocessing strategies?
- Training and testing error rates for $K = 1, 10$ and 100 .

Q5. Survey (4%)

Please give an estimate of how much time you spent on this assignment. Note that you will not be given a higher or lower grade if you spend a lot of time or very little time. We just want an honest estimate. Other feedback are welcomed. Please note that regardless of positive or negative feedback, your grades won’t be affected of course. In other words, as long as your report does not leave this section blank, you will get the full 4%.